Self-Supervised Multi-Modal World Model *with* 4D Space-Time Embedding

Lance Legel *,[1,2]   Qin Huang [3]   Brandon Voelker [4]   Daniel Neamati [5]   Jeff Rose [6]   Patrick Alan Johnson [7]
Favyen Bastani [7]   James Hennessy [8]   Robert Guralnick [2]   Douglas Soltis [2]   Pamela Soltis [2]   Shaowen Wang [9,10]

[1] Ecological Intelligence, Inc.   [2] University of Florida   [3] Arizona State University   [4] University of Houston   [5] Stanford University   [6] SpatialLogic.com
[7] Allen Institute for AI   [8] Georgia Institute of Technology   [9] University of Illinois Urbana-Champaign   [10] NSF I-GUIDE   * lance@ecodash.ai

DeepEarth Architecture

Figure 1. Self-supervised multi-modal world model for planetary science, simulation, & planning.

- **Earth4D** encodes (x, y, z, t) coordinates into learnable space-time embeddings
- DeepEarth fuses multi-modal (*e.g.* vision-language) and space-time embeddings
- Model trains by masked reconstruction of multi-modal data across space-time

Key Contributions

1. **Space-Time Positional Encoder:** Unify all kinds of physical data modalities
2. **Multi-Scale 4D Geospatial Simulator:** Bridge planetary-to-cellular dynamics
3. **4D Learned Hash Probing:** Differentiably map $(x,y,z,t) \rightarrow$ embedding indices

State-of-the-Art Ecological Prediction Benchmark

Live Fuel Moisture Content (LFMC) measures vegetation water for wildfire risk.
Training: Earth4D encodes (x, y, z, t) coordinates of LFMC metrics to 192D, concatenates with learnable plant species embeddings, MLP predicts LFMC.
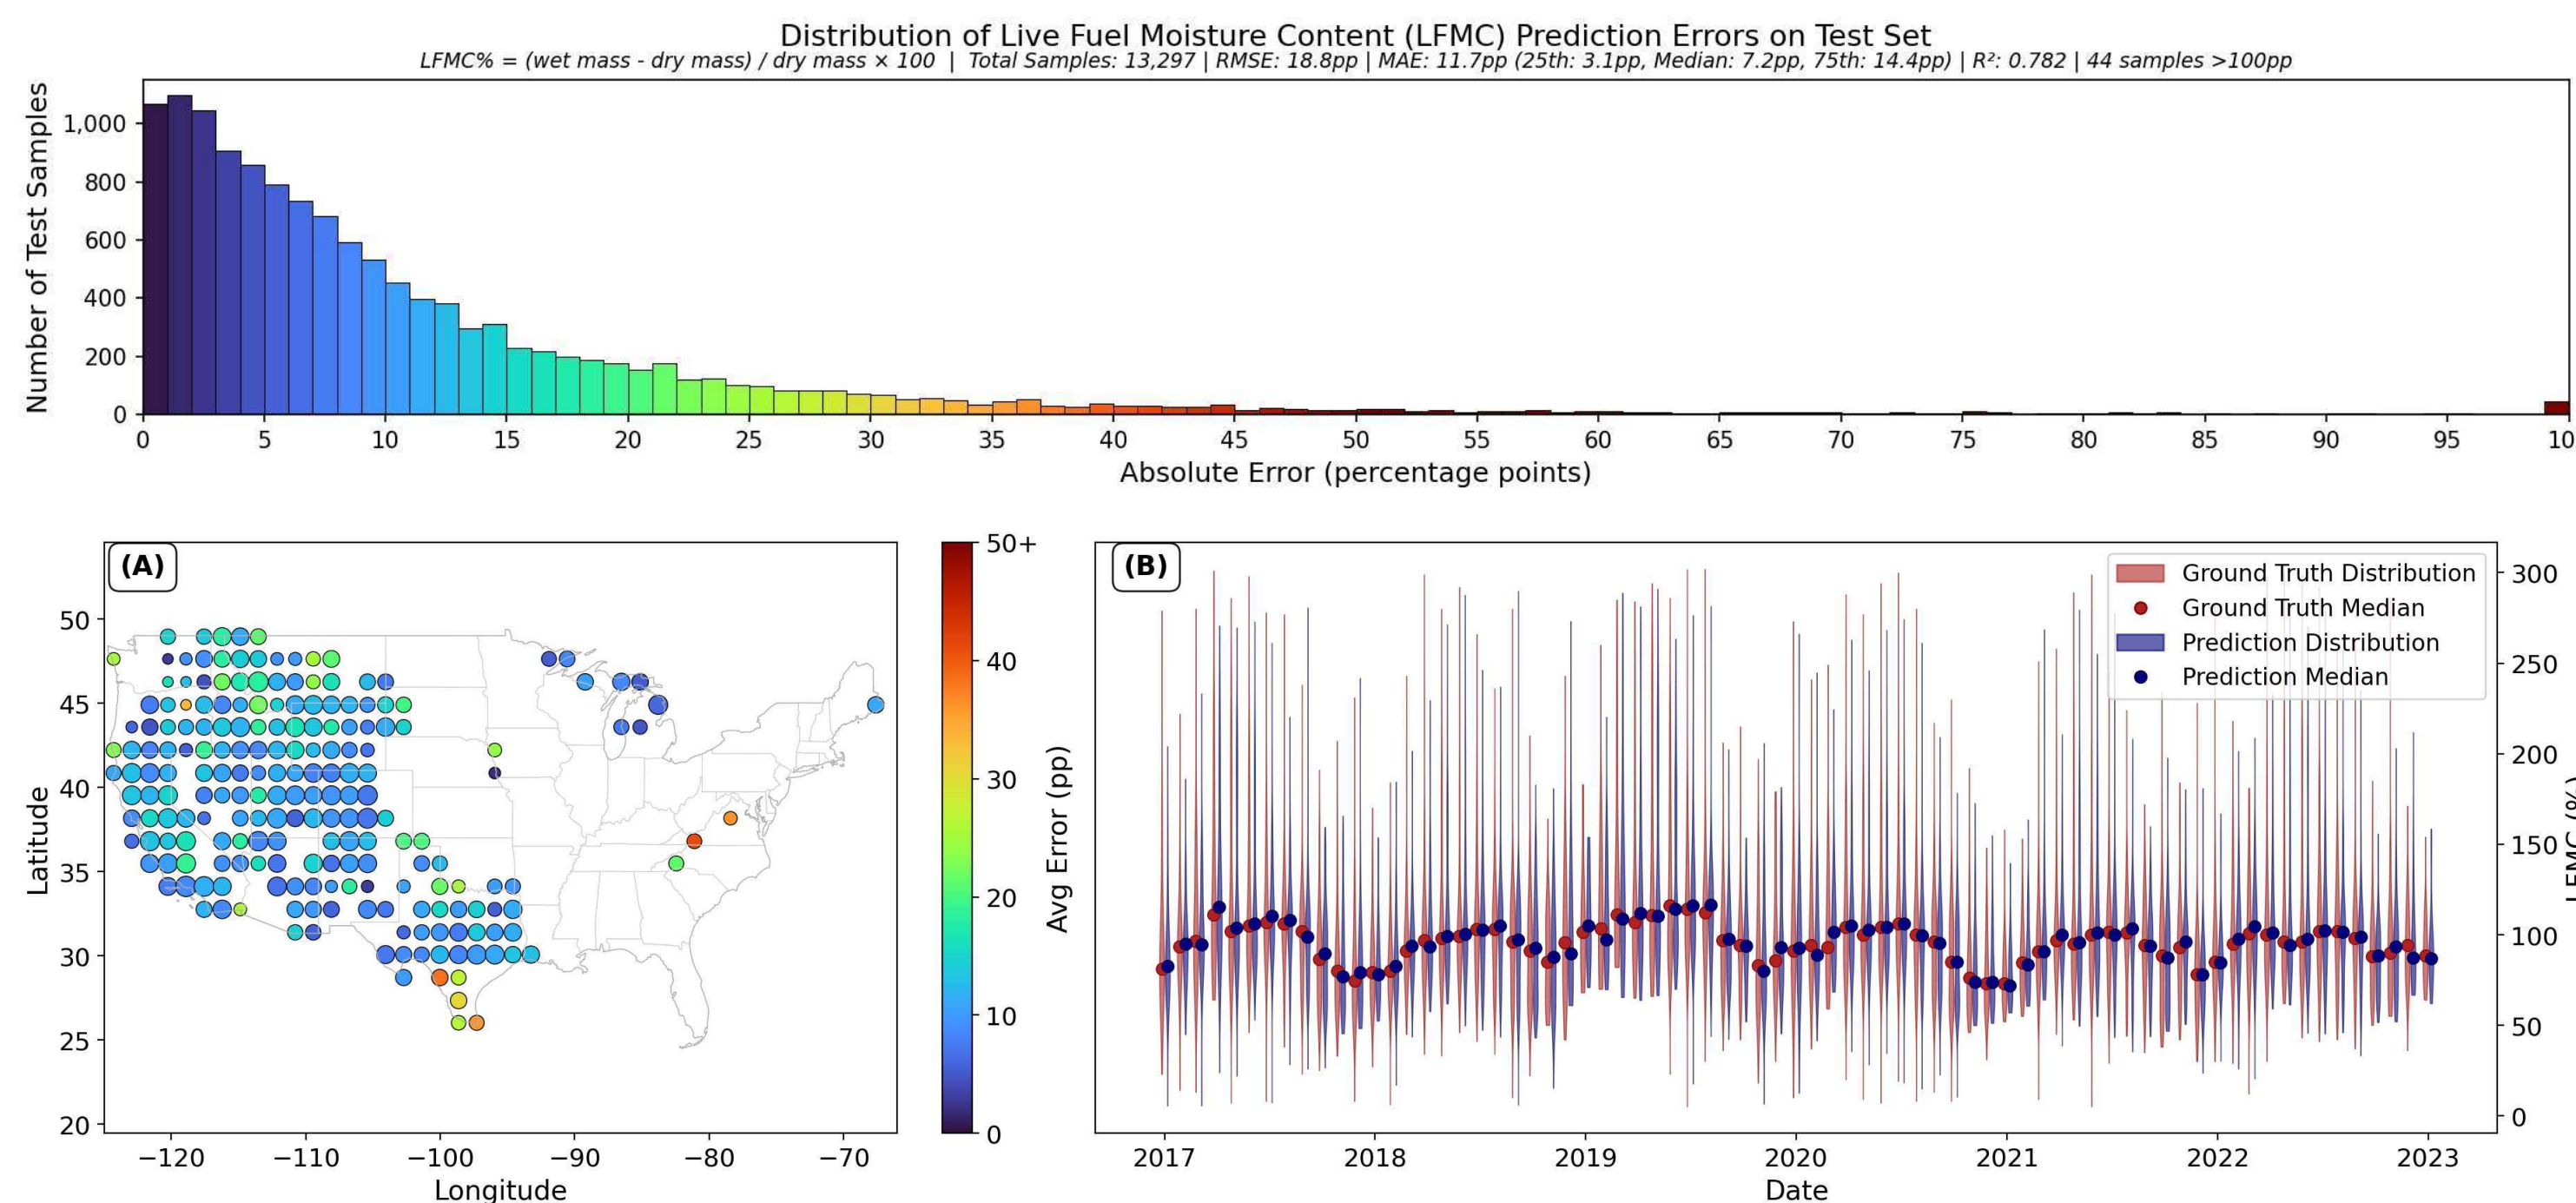**Result:** Earth4D outperforms pre-trained Vision Transformer with less input data.

Figure 4. Earth4D LFMC test predictions on Allen Institute for AI (AI2) Globe-LFMC 2.0 benchmark.

Model	Data Inputs	MAE	RMSE	R²
AI2 (Pre-trained ViT)	(x,y,z,t) + Species + Vision (Remote Sensing)	12.6pp	18.9pp	0.72
Earth4D	(x,y,z,t) + Species	**11.7pp**	**18.7pp**	**0.783**

Key Finding: Earth4D surpasses AI2's foundation model without satellite, weather, or topography data. Only coordinates and species names are used.
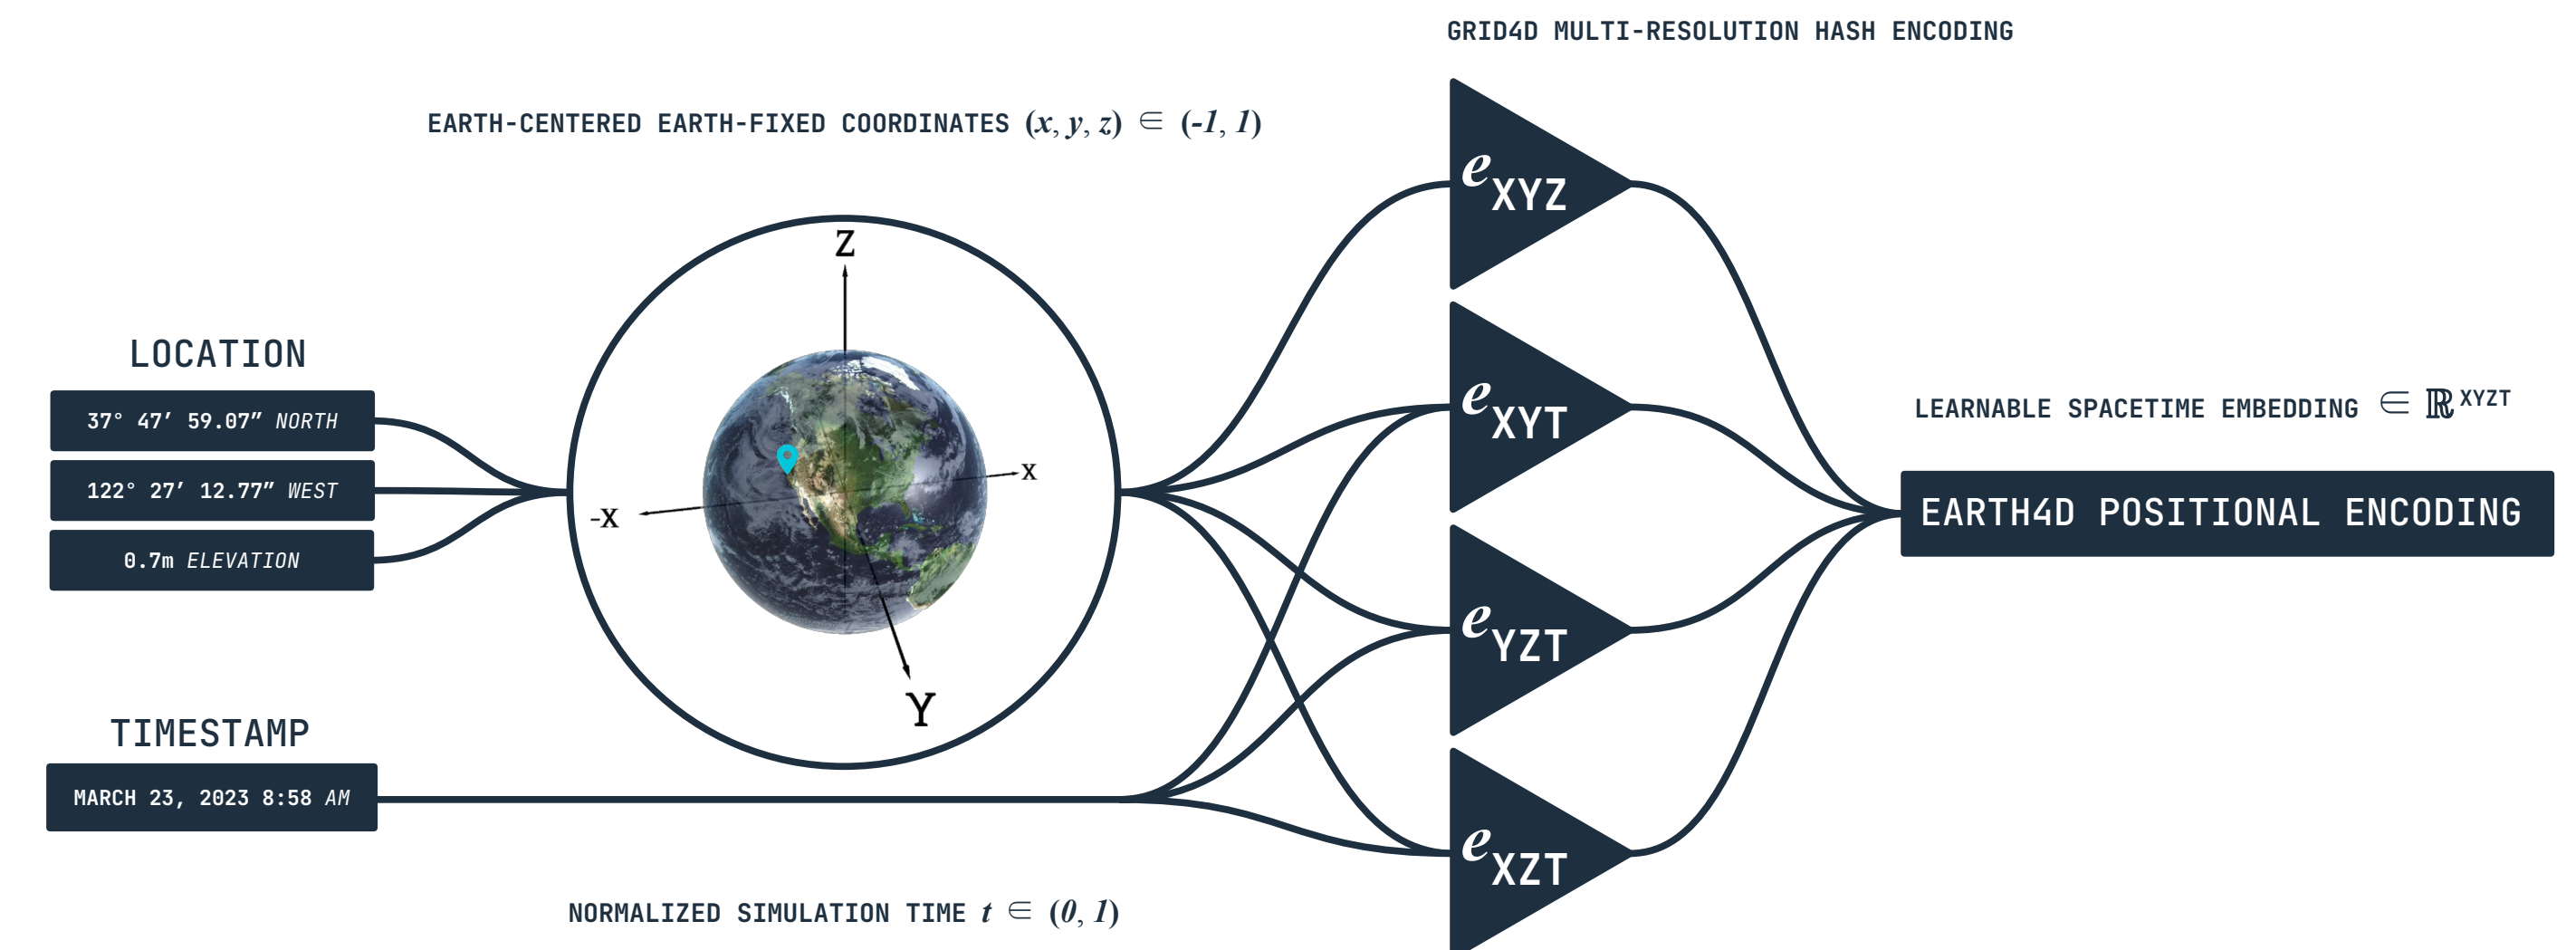
Earth4D Space-Time Positional Encoding

Figure 2. Multi-resolution hash encoding extended to 4D space-time across the planet over years.

Earth4D: **multi-resolution hash encoder** with four parallel 3D grids.
Spatial (xyz): static structure. **Spatio-temporal** (xyt, yzt, xzt): dynamics.
Geographic: Maps $(latitude, longitude, elevation, time) \rightarrow (x, y, z, t)$

## Joint Embeddings Across Spatio-Temporal Scales

Figure 3. 24 resolution levels per grid (xyz, xyt, yzt, xzt), up to 2^{22} entries each.
Output: 192D trainable embedding per (x,y,z,t) coordinate.

Learned Hash Probing

Hash encoding compresses features into fixed memory, but collisions hurt accuracy.
Learned hash probing optimizes memory allocation end-to-end.
Performance: 33% fewer collisions; MAE improved 27%; R² improved 30%.

Code Demo: Space-Time Positional Encoding

```
# https://github.com/legel/deepearth
from deepearth.encoders.xyzt.earth4d import Earth4D
world_model = Earth4D()
embeddings = world_model(
    # Bletchley Park (Turing breaks Enigma, 1941)
    (51.9976, -0.7416, 110, "1941-06-01 09:00 GMT"),
    # Carnegie Mellon (Hinton invents Boltzmann Machines, 1985)
    (40.4433, -79.9436, 270, "1985-01-15 10:00 ET"),
    # CERN (Berners-Lee invents WWW, 1989)
    (46.2330, 6.0557, 430, "1989-03-12 10:00 CET"),
    # Mila, Quebec (World Modeling Workshop 2026)
    (45.5308, -73.6128, 63, "2026-02-04 11:00 ET"),
)
# embeddings.shape: [4, 192] -- trainable space-time features
```

References

1. Müller et al. "Instant Neural Graphics Primitives with a Multiresolution Hash Encoding." ACM SIGGRAPH, 2022.
2. Takikawa et al. "Compact Neural Graphics Primitives with Learned Hash Probing." SIGGRAPH Asia, 2023.
3. Xu et al. "Grid4D: 4D Decomposed Hash Encoding for High-Fidelity Dynamic Gaussian Splatting." NeurIPS, 2024.
4. Yebra et al. "Globe-LFMC 2.0, an enhanced and updated dataset for live fuel moisture content research." Scientific Data, 2024.
5. Tseng et al. "Galileo: Learning Global and Local Features of Many Remote Sensing Modalities." ICML, 2025.
6. Johnson et al. "High-Resolution LFMC Maps for Wildfire Risk From Multimodal Earth Observation Data." PMLR, 2025.